

1. What is the estimated depth of a Decision Tree trained (unrestricted) on a one million instance training set?

Ans:- The estimated depth of a decision tree trained on a one million instance training set would depend on various factors such as the complexity of the data, the number of features, and the decision tree algorithm's specific implementation. However, an unrestricted decision tree trained on a large dataset like that could potentially have a considerable depth.

In general, decision trees have the tendency to grow deeper when trained on larger datasets as they aim to capture more intricate patterns and relationships within the data. The depth of a decision tree represents the number of splits or levels it has, with each split representing a decision based on a feature or attribute.

To provide a rough estimate, decision trees trained on large datasets like the one described (one million instances) can have depths ranging from a few dozen levels to several hundred levels or even more, depending on the complexity of the data and the decision tree algorithm used.

It's worth noting that extremely deep decision trees may lead to overfitting, where the tree becomes overly specific to the training data and performs poorly on unseen data. Therefore, in practice, it is common to apply pruning techniques or use other strategies to limit the depth and complexity of decision trees and avoid overfitting.

1. Is the Gini impurity of a node usually lower or higher than that of its parent? Is it always lower/greater, or is it usually lower/greater?

Ans:- The Gini impurity of a node in a decision tree is typically lower than or equal to the Gini impurity of its parent node. However, it is not always guaranteed to be strictly lower or strictly greater. The Gini impurity measures the level of impurity or uncertainty in a node, with lower values indicating a higher degree of purity.

During the construction of a decision tree, the splitting criterion is typically chosen to minimize the impurity of the resulting child nodes. This means that the impurity of a node's children should be lower than or equal to the impurity of the parent node. However, it is possible for the impurity to remain the same after a split if the feature being used for the split does not effectively separate the classes or if the impurity measure does not change significantly.

In practice, the Gini impurity tends to decrease as the tree grows deeper, as each split aims to partition the data into more homogeneous subsets. However, there may be cases where a particular split results in a higher impurity due to the specific characteristics of the data. Overall, the goal of the decision tree algorithm is to iteratively reduce the impurity as much as possible during the construction process.

1. Explain if its a good idea to reduce max depth if a Decision Tree is overfitting the training set?

Ans:- Reducing the maximum depth of a decision tree can be a good idea if the tree is overfitting the training set. Overfitting occurs when the decision tree captures the noise or random variations in the training data, leading to poor generalization on unseen data. By reducing the maximum depth, the decision tree becomes simpler and less likely to capture complex and noisy patterns in the training data.

Here are a few reasons why reducing the maximum depth can help address overfitting:

Reduced model complexity: Decision trees with a smaller depth have fewer levels or branches, resulting in a simpler model. A simpler model is less prone to overfitting as it focuses on capturing the most important and generalizable patterns in the data.

Less sensitive to noise: Decision trees with a smaller depth are less likely to capture noise or outliers in the training data. By ignoring the finer details or random fluctuations, the tree becomes more robust and generalizes better to unseen data.

Improved generalization: A shallower decision tree tends to have a higher bias and lower variance. This bias-variance trade-off can help improve the model's generalization performance by reducing the chances of overfitting.

However, it's important to note that reducing the maximum depth too much can lead to underfitting, where the decision tree becomes too simple and fails to capture important patterns in the data. It's essential to find the right balance between model complexity and generalization by tuning the hyperparameters, such as the maximum depth, based on cross-validation or other evaluation metrics.

In summary, reducing the maximum depth of a decision tree can be an effective strategy to mitigate overfitting by promoting simplicity, reducing sensitivity to noise, and improving generalization.

1. Explain if its a good idea to try scaling the input features if a Decision Tree underfits the training set?

Ans:- Scaling the input features is generally not necessary or beneficial for decision trees because they are not affected by the scale of the features. Decision trees make splits based on the values of individual features, and the scale of the features does not impact the relative ordering or relationships between the features.

Here are a few points to consider:

Invariant to monotonic transformations: Decision trees are invariant to monotonic transformations of the input features. Scaling the features, such as standardizing them or normalizing them, involves a monotonic transformation that preserves the relative order of values within each feature. As a result, scaling does not affect the decision boundaries or the structure of the decision tree.

Splitting based on threshold values: Decision trees split the data based on threshold values of individual features. Scaling the features does not change the threshold values or the order in which the splits are made. Therefore, the decision tree will produce the same structure and

predictions whether the features are scaled or not.

Focus on feature relationships: Decision trees are more concerned with the relationships and interactions between features rather than their absolute values or scales. They can capture nonlinear relationships and interactions without the need for feature scaling.

However, there may be some exceptional cases where scaling could be relevant for decision trees. For example, if there are features with significantly different scales, and the algorithm used for feature selection or pre-processing requires scaling (e.g., feature importance calculations based on scaling), then scaling might be considered. Additionally, if there are other algorithms or models in the pipeline that require scaled features, it might be practical to scale the features consistently.

In summary, scaling the input features is generally unnecessary for decision trees as they are not affected by the scale of the features. Decision trees make splits based on the relative order of feature values, and scaling does not impact these relationships or the resulting decision boundaries.

1. How much time will it take to train another Decision Tree on a training set of 10 million instances if it takes an hour to train a Decision Tree on a training set with 1 million instances?

Ans:- The time it takes to train a decision tree on a larger training set is not necessarily linearly proportional to the size of the training set. The training time can be influenced by various factors such as the complexity of the data, the specific implementation of the decision tree algorithm, the hardware used, and any optimizations applied.

However, as a rough estimate, we can assume that training time could be approximately linearly proportional to the number of instances in the training set. In this case, if it takes an hour to train a decision tree on a training set with 1 million instances, we can estimate that training a decision tree on a training set with 10 million instances could take around 10 hours.

It's important to note that this is just a rough estimate and the actual training time can vary based on the factors mentioned above. Additionally, if you are using parallel processing or distributed computing techniques, it might be possible to reduce the training time significantly.

1. Will setting `presort=True` speed up training if your training set has 100,000 instances?

Ans:- The `presort` parameter in scikit-learn's decision tree algorithm determines whether to presort the data for faster tree building. By setting `presort=True`, the algorithm will pre-sort the data based on the feature values, which can speed up training in some cases.

However, it's important to note that enabling `presort=True` comes with a trade-off. Pre-sorting the data requires additional computational time and memory, which can be significant for large datasets. The benefit of using `presort=True` typically diminishes as the dataset size increases. Therefore, enabling `presort=True` is generally recommended for smaller datasets or when the algorithm's performance with `presort=True` is better than without it.

For a training set with 100,000 instances, enabling `presort=True` may or may not provide a noticeable speedup. It depends on the specific characteristics of the dataset, such as the number of features, the complexity of the data, and the available computational resources. It's recommended to try training with and without `presort=True` and measure the training time to determine the impact in your specific scenario.

1. Follow these steps to train and fine-tune a Decision Tree for the moons dataset:

a. To build a moons dataset, use `make_moons(n_samples=10000, noise=0.4)`.

Sol: `from sklearn.datasets import make_moons`

`X, y = make_moons(n_samples=10000, noise=0.4)`

b. Divide the dataset into a training and a test collection with `train test split()`.

Sol: `from sklearn.model_selection import train_test_split`

`X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)`

c. To find good hyperparameters values for a `DecisionTreeClassifier`, use grid search with cross-validation (with the `GridSearchCV` class). Try different values for max leaf nodes.

Sol: `from sklearn.tree import DecisionTreeClassifier from sklearn.model_selection import GridSearchCV`

`param_grid = {'max_leaf_nodes': [None, 5, 10, 20, 50]} tree_clf = DecisionTreeClassifier(random_state=42) grid_search = GridSearchCV(tree_clf, param_grid, cv=5) grid_search.fit(X_train, y_train)`

`print("Best hyperparameters:", grid_search.best_params)`

d. Use these hyperparameters to train the model on the entire training set, and then assess its output on the test set. You can achieve an accuracy of 85 to 87 percent.

Sol: `best_tree_clf = DecisionTreeClassifier(max_leaf_nodes=grid_search.best_params['max_leaf_nodes'], random_state=42)`  
`best_tree_clf.fit(X_train, y_train)`

```
accuracy = best_tree_clf.score(X_test, y_test) print("Accuracy on the test set:", accuracy)
```

1. Follow these steps to grow a forest:

a. Using the same method as before, create 1,000 subsets of the training set, each containing 100 instances chosen at random. You can do this with Scikit-ShuffleSplit Learn's class.

Sol: from sklearn.model\_selection import ShuffleSplit

```
n_subsets = 1000 subset_size = 100 shuffle_split = ShuffleSplit(n_splits=n_subsets, test_size=subset_size, random_state=42)
```

```
subsets = [] for train_index, in shuffle_split.split(X_train): subset = X_train[train_index] subsets.append(subset)
```

b. Using the best hyperparameter values found in the previous exercise, train one Decision Tree on each subset. On the test collection, evaluate these 1,000 Decision Trees. These Decision Trees would likely perform worse than the first Decision Tree, achieving only around 80% accuracy, since they were trained on smaller sets.

Sol: from sklearn.tree import DecisionTreeClassifier

```
trees = [] for subset in subsets: tree = DecisionTreeClassifier(max_leaf_nodes=grid_search.best_params['max_leaf_nodes'], random_state=42) tree.fit(subset, y_train[:subset_size]) trees.append(tree)
```

c. Now the magic begins. Create 1,000 Decision Tree predictions for each test set case, and keep only the most common prediction (you can do this with SciPy's mode() function). Over the test collection, this method gives you majority-vote predictions.

Sol: import numpy as np from scipy.stats import mode

```
predictions = np.empty([n_subsets, len(X_test)]) for i, tree in enumerate(trees): predictions[i] = tree.predict(X_test)
```

```
ensemble_predictions, = mode(predictions, axis=0)
```

d. On the test range, evaluate these predictions: you should achieve a slightly higher accuracy than the first model (approx 0.5 to 1.5 percent higher). You've successfully learned a Random Forest classifier!

Sol: ensemble\_accuracy = np.sum(ensemble\_predictions == y\_test) / len(y\_test) print("Accuracy of the random forest on the test set:", ensemble\_accuracy)